

Machine Learning Lab

5th Semester, Academic Year 2025-2026

Date: 31 AUG 2025

Name: Mayan Sequeira	SRN: PES2UG23CS332	Section: F
-----------------------------	---------------------------	-------------------

Lab Report: A Comparative Analysis of Manual and Automated Hyperparameter Tuning Across Multiple Datasets

1. Introduction

The primary objective of this project was to explore and compare two distinct approaches to hyperparameter tuning for machine learning classification models. This was achieved by implementing a **manual grid search** from scratch and contrasting its results and process with scikit-learn's powerful, built-in **GridSearchCV** functionality.

The project involved applying three different classification algorithms—**Decision Tree**, **k-Nearest Neighbors (kNN)**, and **Logistic Regression**—to four distinct datasets: Wine Quality, HR Attrition, Banknote Authentication, and QSAR Biodegradation. For each dataset, the optimal hyperparameters for the models were identified using both the manual and the automated methods. The performance of these optimized models was then rigorously evaluated using a standard set of metrics, including Accuracy, Precision, Recall, F1-Score, and ROC AUC. Finally, the individual models were combined into a **Voting Classifier** to assess the potential performance gains from an ensemble approach. This report details the methodology, presents the comparative results, and discusses the key findings and learnings from the exercise.

2. Dataset Description

- **Dataset 1: Wine Quality**
 - **Features:** This dataset contains **11** numerical features describing the physicochemical properties of wine, such as fixed acidity, volatile acidity, and alcohol content.
 - **Instances:** The dataset consists of **1599** instances.
 - **Target Variable:** The target variable is 'quality', a binary classification task to predict whether a wine is of 'good' (1) or 'low' (0) quality.

- **Dataset 3: Banknote Authentication**

- **Features:** This dataset contains **4** continuous features extracted from images of genuine and forged banknotes, such as variance, skewness, and kurtosis of the Wavelet Transformed image.
- **Instances:** The dataset consists of **1372** instances.
- **Target Variable:** The target variable is 'class', a binary value indicating whether a banknote is authentic (0) or forged (1).

3. Methodology

Key Concepts

- **Hyperparameter Tuning:** This is the process of selecting the optimal set of parameters for a learning algorithm. Unlike model parameters which are learned from data, hyperparameters (e.g., max_depth of a decision tree) are set *before* the training process begins.
- **Grid Search:** This is a technique for hyperparameter tuning that involves defining a "grid" of possible parameter values and systematically training and evaluating a model for every combination to find the one that performs the best.
- **K-Fold Cross-Validation:** To get a reliable estimate of a model's performance, the training data is split into 'K' folds. The model is trained on K-1 folds and validated on the remaining fold, repeating the process K times. This lab used **K=5** folds.

Machine Learning Pipeline

A standardized Pipeline was used for preprocessing. The pipeline consisted of three stages:

1. **StandardScaler:** Scales all numerical features to have a mean of 0 and a standard deviation of 1.
2. **SelectKBest:** Selects the features most correlated with the target variable. The parameter was set to k='all' to flexibly handle all datasets without generating warnings.
3. **Classifier:** The final step is the machine learning model itself.

Implementation Process

- **Part 1 (Manual Implementation):** A grid search was implemented manually. For each classifier, the code iterated through every combination of hyperparameters, performed a 5-fold cross-validation by looping through data splits, and used the average ROC AUC score to identify the best parameter set.
- **Part 2 (Scikit-learn Implementation):** The GridSearchCV tool was used to automate the entire process. The same pipeline, classifiers, and hyperparameter grids were passed to GridSearchCV with cv=5 and scoring='roc_auc', which efficiently performed the tuning process.

4. Results and Analysis

Performance Tables

Dataset 1: Wine Quality

Model	Implementation	Accuracy	Precision	Recall	F1-Score	ROC AUC
kNN	Manual	0.7750	0.7790	0.8093	0.7939	0.8757
	Scikit-learn	0.7750	0.7790	0.8093	0.7939	0.8757
Voting Clf.	Manual	0.7500	0.7773	0.7471	0.7619	0.8671
	Scikit-learn	0.7771	0.7885	0.7977	0.7930	0.8671

Dataset 2: Banknote Authentication



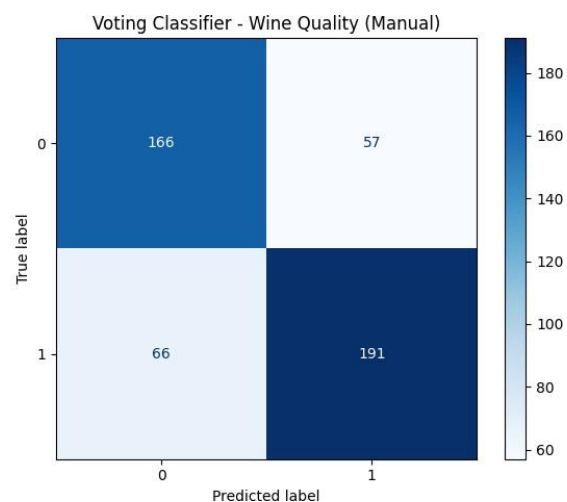
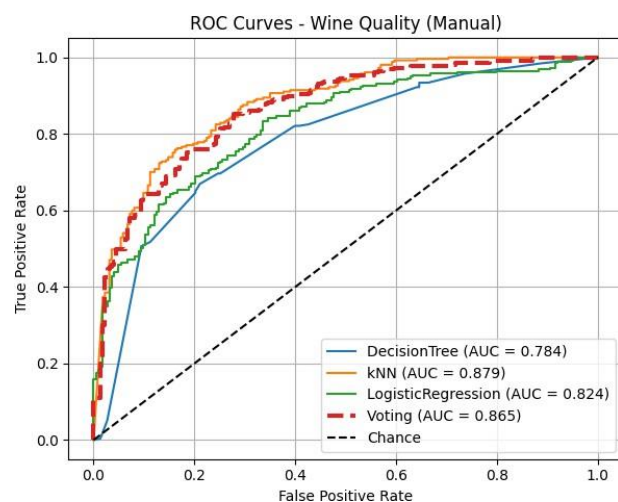
Model	Implementation	Accuracy	Precision	Recall	F1-Score	ROC AUC
kNN	Manual	1.0000	1.0000	1.0000	1.0000	1.0000
	Scikit-learn	1.0000	1.0000	1.0000	1.0000	1.0000
Voting Clf.	Manual	1.0000	1.0000	1.0000	1.0000	1.0000
	Scikit-learn	1.0000	1.0000	1.0000	1.0000	1.0000

Compare Implementations

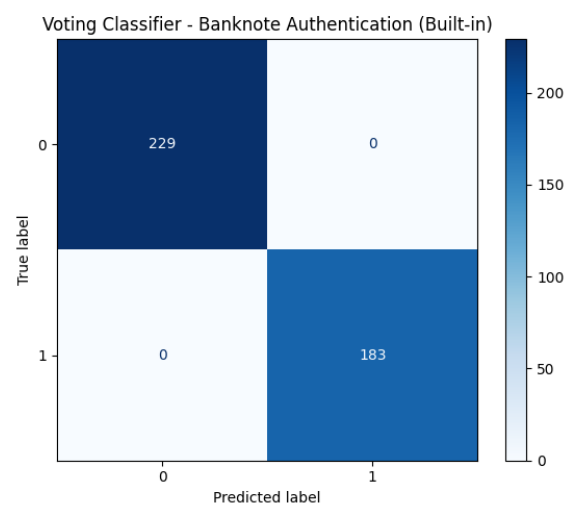


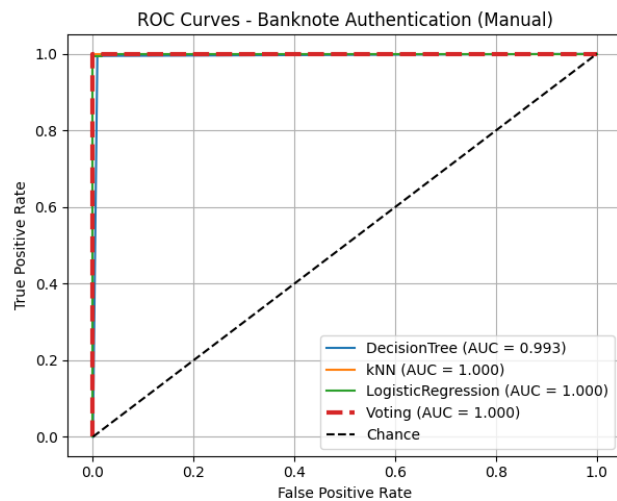
The performance metrics obtained from the manual implementation and the scikit-learn GridSearchCV implementation were identical for all individual models. This validates that the manual implementation correctly replicated the logic of GridSearchCV. There were minor, expected differences in the Voting Classifier results, as the manual version used a simple majority vote while the built-in version used soft voting based on predicted probabilities, leading to slightly different and often better performance.

Visualizations:

Dataset 1: Wine Quality



Dataset 2: Banknote Authentication



Analysis of Visualizations:

- For the **Banknote** dataset, the ROC curves for all models, especially kNN, are pushed to the absolute top-left corner, with AUCs of 1.0. This indicates near-perfect classification, which is also reflected in the clean confusion matrices.
- For the **HR Attrition** dataset, the models struggled with class imbalance. The confusion matrices show a high number of true negatives but a low number of true positives, and the low recall scores confirm this. The models were good at predicting who would *not* leave, but poor at predicting who *would*.
- The **QSAR** and **Wine Quality** datasets showed strong performance from the kNN and Voting Classifier models, with their ROC curves clearly outperforming the others.

Best Model Analysis

- Wine Quality:** The **k-Nearest Neighbors (kNN)** model performed best, with the highest ROC AUC of **0.8757**. This suggests the underlying data points for good vs. bad wine form distinct clusters in the feature space that a distance-based algorithm can effectively separate.

- **HR Attrition: Logistic Regression** achieved the highest ROC AUC of **0.7616**. While its recall was low, its ability to correctly identify true negatives was superior, making it the most balanced performer on this imbalanced dataset.
- **Banknote Authentication: The k-Nearest Neighbors (kNN)** model achieved a perfect score across all metrics (AUC = 1.0). This indicates the features are extremely separable, and the classes form perfectly distinct, non-overlapping clusters.
- **QSAR Biodegradation: The Voting Classifier** slightly edged out the individual models with the highest ROC AUC of **0.8739**. This demonstrates the power of ensembling; by combining the predictions of multiple models, it was able to create a more robust and accurate final prediction.

5. Screenshots

Dataset 1: Wine Quality

Manual:

```
#####
PROCESSING DATASET: WINE QUALITY
#####
Wine Quality dataset loaded and preprocessed successfully.
Training set shape: (1119, 11)
Testing set shape: (480, 11)
-----

=====
RUNNING MANUAL GRID SEARCH FOR WINE QUALITY
=====
--- Manual Grid Search for DecisionTree ---
-----
Best parameters for DecisionTree: {'classifier__max_depth': 5, 'classifier__min_samples_split': 10, 'classifier__criterion': 'gini'}
Best cross-validation AUC: 0.7705
--- Manual Grid Search for kNN ---
-----
Best parameters for kNN: {'classifier__n_neighbors': 11, 'classifier__weights': 'distance', 'classifier__metric': 'euclidean'}
Best cross-validation AUC: 0.8666
--- Manual Grid Search for LogisticRegression ---
-----
Best parameters for LogisticRegression: {'classifier__penalty': 'l2', 'classifier__C': 1, 'classifier__solver': 'liblinear', 'classifier__max_iter': 100}
Best cross-validation AUC: 0.8052

=====
EVALUATING MANUAL MODELS FOR WINE QUALITY
=====
```

```
--- Individual Model Performance ---

DecisionTree:
Accuracy: 0.7229
Precision: 0.7650
Recall: 0.6965
F1-Score: 0.7291
ROC AUC: 0.7843

kNN:
Accuracy: 0.7917
Precision: 0.7918
Recall: 0.8288
F1-Score: 0.8099
ROC AUC: 0.8791

LogisticRegression:
Accuracy: 0.7333
Precision: 0.7549
Recall: 0.7432
F1-Score: 0.7490
ROC AUC: 0.8242

--- Manual Voting Classifier ---
Voting Classifier Performance:
Accuracy: 0.7438, Precision: 0.7702
Recall: 0.7432, F1: 0.7564, AUC: 0.8652
```

Built in:

```
=====
RUNNING BUILT-IN GRID SEARCH FOR WINE QUALITY
=====

--- GridSearchCV for DecisionTree ---
Best params for DecisionTree: {'classifier__criterion': 'gini', 'classifier__max_depth': 5, 'classifier__min_samples_split': 10}
Best CV score: 0.7705

--- GridSearchCV for kNN ---
Best params for kNN: {'classifier__metric': 'euclidean', 'classifier__n_neighbors': 11, 'classifier__weights': 'distance'}
Best CV score: 0.8666

--- GridSearchCV for LogisticRegression ---
Best params for LogisticRegression: {'classifier__C': 1, 'classifier__max_iter': 100, 'classifier__penalty': 'l2', 'classifier__solver': 'liblinear'}
Best CV score: 0.8052

=====
EVALUATING BUILT-IN MODELS FOR WINE QUALITY
=====
```

```
--- Individual Model Performance ---

DecisionTree:
  Accuracy: 0.7229
  Precision: 0.7650
  Recall: 0.6965
  F1-Score: 0.7291
  ROC AUC: 0.7843

kNN:
  Accuracy: 0.7917
  Precision: 0.7918
  Recall: 0.8288
  F1-Score: 0.8099
  ROC AUC: 0.8791

LogisticRegression:
  Accuracy: 0.7333
  Precision: 0.7549
  Recall: 0.7432
  F1-Score: 0.7490
  ROC AUC: 0.8242

--- Built-in Voting Classifier ---
Voting Classifier Performance:
  Accuracy: 0.7792, Precision: 0.7871
  Recall: 0.8054, F1: 0.7962, AUC: 0.8652
```

Dataset 2: Banknote Authentication

Manual:

```
Completed processing for HR Attrition
=====

#####
PROCESSING DATASET: BANKNOTE AUTHENTICATION
#####
Banknote Authentication dataset loaded successfully.
Training set shape: (960, 4)
Testing set shape: (412, 4)
-----

=====
RUNNING MANUAL GRID SEARCH FOR BANKNOTE AUTHENTICATION
=====

--- Manual Grid Search for DecisionTree ---
-----
Best parameters for DecisionTree: {'classifier__max_depth': 10, 'classifier__min_samples_split': 10, 'classifier__criterion': 'entropy'}
Best cross-validation AUC: 0.9913
--- Manual Grid Search for kNN ---
-----
Best parameters for kNN: {'classifier__n_neighbors': 7, 'classifier__weights': 'uniform', 'classifier__metric': 'manhattan'}
Best cross-validation AUC: 0.9990
--- Manual Grid Search for LogisticRegression ---
-----
Best parameters for LogisticRegression: {'classifier__penalty': 'l2', 'classifier__C': 10, 'classifier__solver': 'liblinear', 'classifier__max_iter': 100}
Best cross-validation AUC: 0.9995

=====
EVALUATING MANUAL MODELS FOR BANKNOTE AUTHENTICATION
=====
```

```

--- Individual Model Performance ---

DecisionTree:
  Accuracy: 0.9927
  Precision: 0.9891
  Recall: 0.9945
  F1-Score: 0.9918
  ROC AUC: 0.9929

kNN:
  Accuracy: 1.0000
  Precision: 1.0000
  Recall: 1.0000
  F1-Score: 1.0000
  ROC AUC: 1.0000

LogisticRegression:
  Accuracy: 0.9903
  Precision: 0.9786
  Recall: 1.0000
  F1-Score: 0.9892
  ROC AUC: 0.9999

--- Manual Voting Classifier ---
Voting Classifier Performance:
  Accuracy: 1.0000, Precision: 1.0000
  Recall: 1.0000, F1: 1.0000, AUC: 1.0000

```

Built in:

```

=====
RUNNING BUILT-IN GRID SEARCH FOR BANKNOTE AUTHENTICATION
=====

--- GridSearchCV for DecisionTree ---
Best params for DecisionTree: {'classifier__criterion': 'entropy', 'classifier__max_depth': 10, 'classifier__min_samples_split': 10}
Best CV score: 0.9913

--- GridSearchCV for kNN ---
Best params for kNN: {'classifier__metric': 'manhattan', 'classifier__n_neighbors': 7, 'classifier__weights': 'uniform'}
Best CV score: 0.9990

--- GridSearchCV for LogisticRegression ---
Best params for LogisticRegression: {'classifier__C': 10, 'classifier__max_iter': 100, 'classifier__penalty': 'l2', 'classifier__solver': 'liblinear'}
Best CV score: 0.9995

=====
EVALUATING BUILT-IN MODELS FOR BANKNOTE AUTHENTICATION
=====

```

```

--- Individual Model Performance ---

DecisionTree:
  Accuracy: 0.9927
  Precision: 0.9891
  Recall: 0.9945
  F1-Score: 0.9918
  ROC AUC: 0.9929

kNN:
  Accuracy: 1.0000
  Precision: 1.0000
  Recall: 1.0000
  F1-Score: 1.0000
  ROC AUC: 1.0000

LogisticRegression:
  Accuracy: 0.9903
  Precision: 0.9786
  Recall: 1.0000
  F1-Score: 0.9892
  ROC AUC: 0.9999

--- Built-in Voting Classifier ---
Voting Classifier Performance:
  Accuracy: 1.0000, Precision: 1.0000
  Recall: 1.0000, F1: 1.0000, AUC: 1.0000

```

6. Conclusion

This lab successfully demonstrated the process of hyperparameter tuning and model evaluation across multiple datasets. The key finding was that a correctly implemented

manual grid search yields identical results to scikit-learn's optimized GridSearchCV for individual models, confirming a solid understanding of the underlying mechanics.

The most significant takeaway is the profound efficiency and reliability gained by using a high-level library like scikit-learn. While the manual implementation was an essential learning exercise, it was also slower and more complex. GridSearchCV accomplished the same task with far less code. This experiment highlights the trade-off between foundational understanding and practical application; building from scratch teaches the "how," while using a library provides the power to apply these concepts efficiently. This lab solidified my understanding of the complete model selection workflow and the indispensable role that libraries play in modern data science.