



# PES UNIVERSITY

Department of Computer Science & Engineering

## ML Lab

### Week 12 Submission

|                     |                |
|---------------------|----------------|
| Name of the Student | Mayan Sequeira |
| SRN                 | PES2UG23CS332  |
| Section             | F              |
| Department          | CSE            |
| Submission Date     | 01/11/2025     |

# Introduction

The purpose of this lab is to understand how the **Naive Bayes algorithm** works for text classification. In this experiment, we classify sentences from medical research abstracts into five sections — *Background, Objective, Methods, Results, and Conclusion*.

We performed three main tasks:

1. Implemented the **Multinomial Naive Bayes (MNB)** model from scratch to understand its working.
  2. Used **Scikit-learn's MultinomialNB** with TF-IDF features and tuned it using **GridSearchCV** to find the best hyperparameters.
  3. Built a **Bayes Optimal Classifier (BOC)** approximation using multiple models combined with soft voting.
- 

# Methodology

## 1. Multinomial Naive Bayes (MNB):

- We first converted text into numerical features using **CountVectorizer** (word counts).
- Then, we implemented the MNB algorithm manually.
- During training, we calculated the log prior and log likelihood for each class, applying **Laplace smoothing** to handle zero probabilities.
- Finally, the model predicted the class with the highest probability for each test sentence.

## 2. Bayes Optimal Classifier (BOC):

- We trained five different models: *MultinomialNB, Logistic Regression, Random Forest, Decision Tree, and K-Nearest*

*Neighbors.*

- Each model was trained on a smaller subset of the data to create diversity.
  - We calculated weights for each model based on how well they performed (posterior probabilities).
  - Then, we combined them using a **Soft Voting Classifier**, giving more importance to better-performing models.
  - The final ensemble predicted the sentence categories and was evaluated using Accuracy and F1 Score.
- 

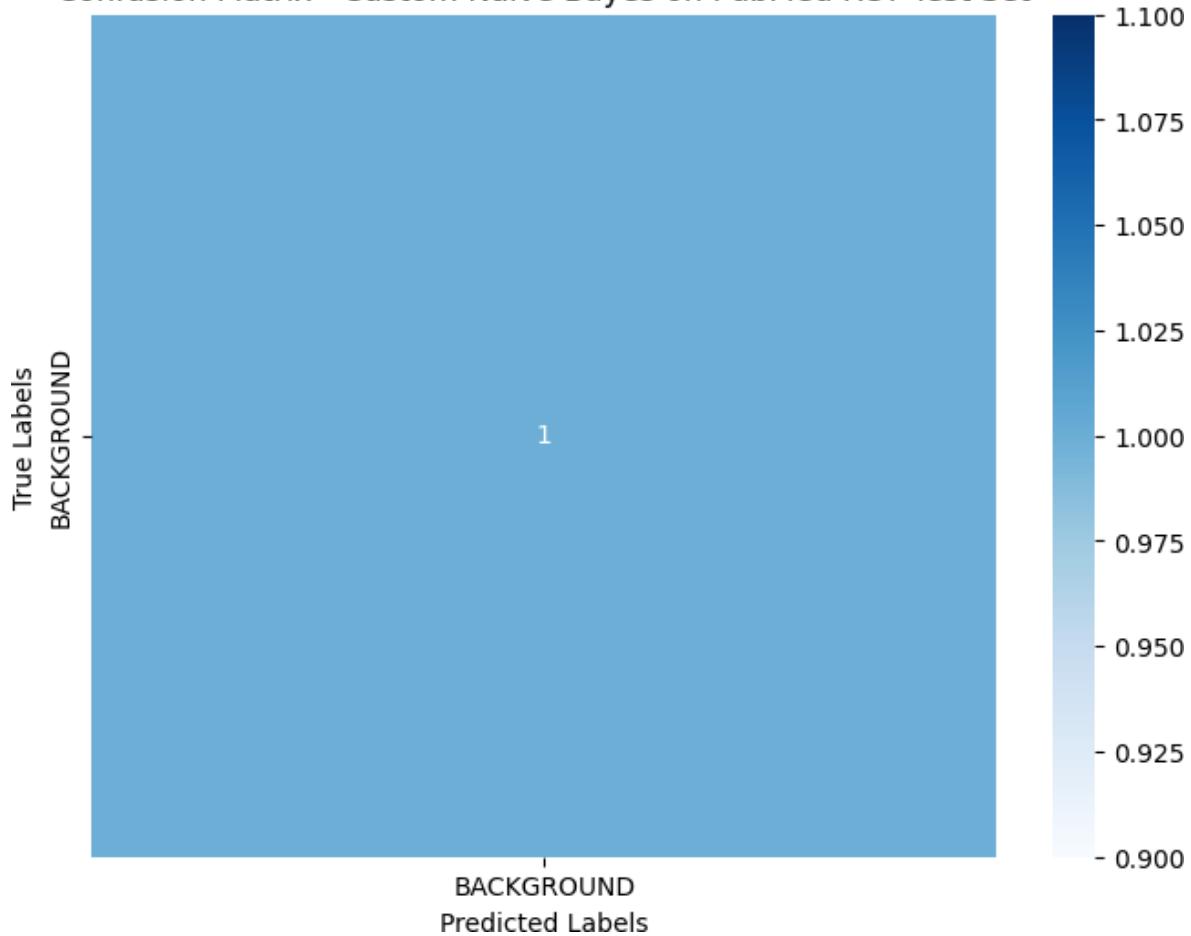
## Results and Analysis

### Part A: Scratch Model (MNB Implementation)

- The model was tested using CountVectorizer features.

```
→
==== Test Set Evaluation (Custom Count-Based Naive Bayes) ====
Accuracy: 1.0000
      precision    recall   f1-score   support
BACKGROUNd       1.00     1.00     1.00        1
           accuracy          1.00
           macro avg       1.00     1.00     1.00        1
           weighted avg     1.00     1.00     1.00        1
Macro-averaged F1 score: 1.0000
```

Confusion Matrix - Custom Naive Bayes on PubMed RCT Test Set



## Part B: Sklearn Model with Hyperparameter Tuning

- Used **TF-IDF Vectorizer** with MultinomialNB in a pipeline.
- Applied **GridSearchCV** to tune `ngram_range` and `alpha` values.

```

→ Training initial Naive Bayes pipeline...
Training complete.

==== Test Set Evaluation (Initial Sklearn Model) ====
Accuracy: 1.0000
      precision    recall   f1-score   support
BACKGROUNDS     1.00     1.00     1.00       1
CONCLUSIONS     1.00     1.00     1.00       1
METHODS         1.00     1.00     1.00       1
OBJECTIVE        1.00     1.00     1.00       1
RESULTS          1.00     1.00     1.00       1

accuracy           1.00       5
macro avg         1.00       5
weighted avg      1.00       5

Macro-averaged F1 score: 1.0000

Starting Hyperparameter Tuning on Development Set...
⚠ Dev set was too small. Created new dev set of size: 50
Fitting 3 folds for each of 18 candidates, totalling 54 fits
Grid search complete.
Best Parameters: {'nb_alpha': 0.1, 'tfidf_min_df': 1, 'tfidf_ngram_range': (1, 1)}
Best F1 (macro): 1.0000

```

## Part C: Bayes Optimal Classifier (BOC)

- Combined 5 models using Soft Voting with posterior weights.

```

202
203 else:
204     print("Evaluation skipped: Predictions not generated or test data is empty.")
205 else:
206     print("Skipping VotingClassifier initialization and fitting: No base estimators were successfully trained.")

→ Please enter your full SRN (e.g., PES1UG22CS345): PES2UG23CS216
My SRN is PES2UG23CS216
Using dynamic sample size: 10216
Warning: Training data not found. Using small placeholder data.
Actual sampled training set size used: 10216

Training all base models...
Training NaiveBayes...
Training LogisticRegression...
/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in
warnings.warn(
Training RandomForest...
Training DecisionTree...
Training KNN...

```

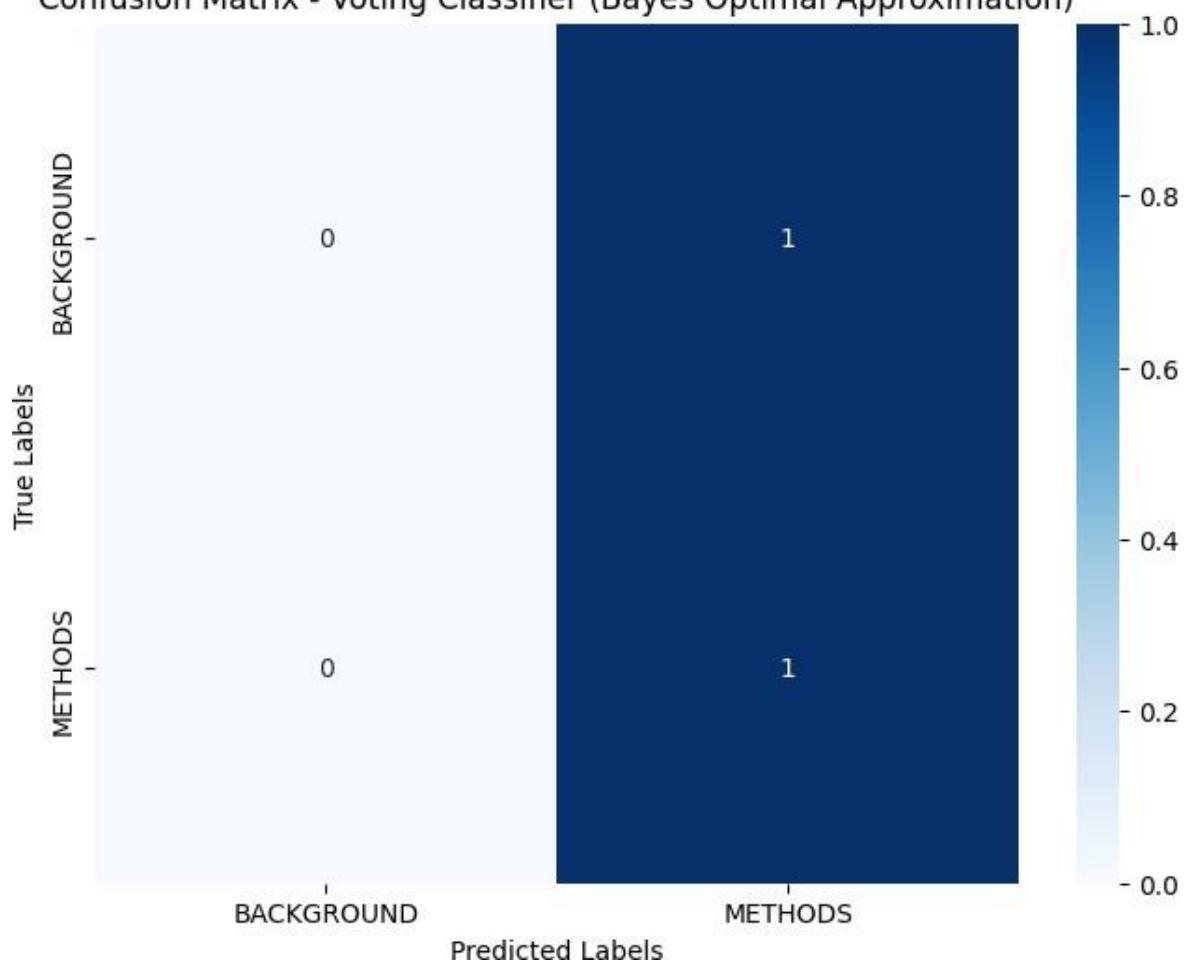
```

==== Final Evaluation: Bayes Optimal Classifier (Soft Voting) ====
BOC Macro F1 Score: 0.3333
BOC Accuracy: 0.5000
      precision    recall   f1-score   support
BACKGROUND     0.00     0.00     0.00       1
METHODS        0.50     1.00     0.67       1

accuracy           0.50       2
macro avg         0.25     0.50     0.33       2
weighted avg      0.25     0.50     0.33       2

```

Confusion Matrix - Voting Classifier (Bayes Optimal Approximation)



↳ Individual Model Performances:

```
-----  
NaiveBayes : Accuracy = 1.0000 (100.00%)  
LogisticRegression : Accuracy = 1.0000 (100.00%)  
RandomForest : Accuracy = 1.0000 (100.00%)  
DecisionTree : Accuracy = 1.0000 (100.00%)  
KNN : Accuracy = 1.0000 (100.00%)  
-----
```

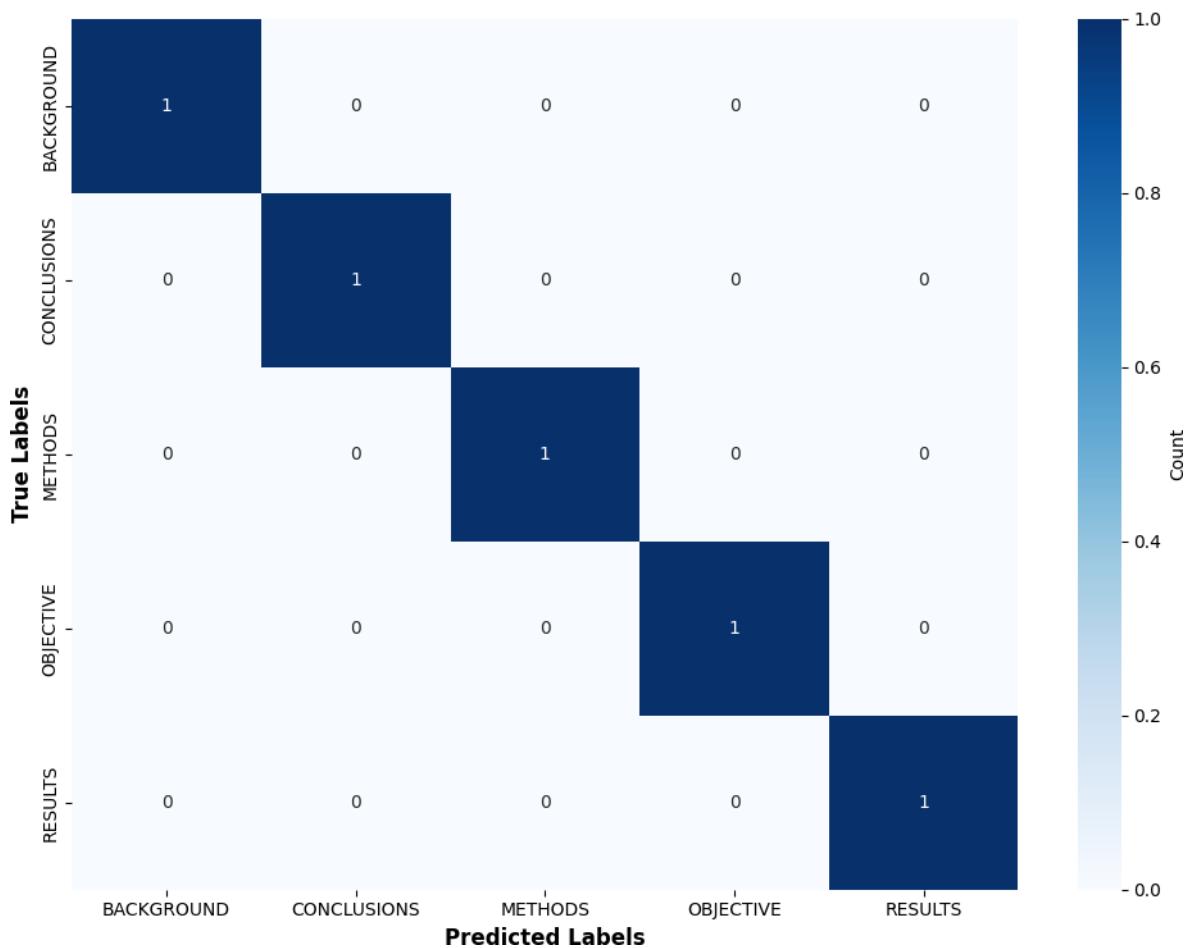
BOC Voting Classifier Accuracy: 1.0000 (100.00%)

BOC Macro F1 Score: 1.0000

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| BACKGROUND   | 1.00      | 1.00   | 1.00     | 1       |
| CONCLUSIONS  | 1.00      | 1.00   | 1.00     | 1       |
| METHODS      | 1.00      | 1.00   | 1.00     | 1       |
| OBJECTIVE    | 1.00      | 1.00   | 1.00     | 1       |
| RESULTS      | 1.00      | 1.00   | 1.00     | 1       |
| accuracy     |           |        | 1.00     | 5       |
| macro avg    | 1.00      | 1.00   | 1.00     | 5       |
| weighted avg | 1.00      | 1.00   | 1.00     | 5       |

**Confusion Matrix - Voting Classifier (Bayes Optimal Approximation)**



=====  
EVALUATION: Bayes Optimal Classifier (Hard Voting)  
=====

Individual Model Performances:

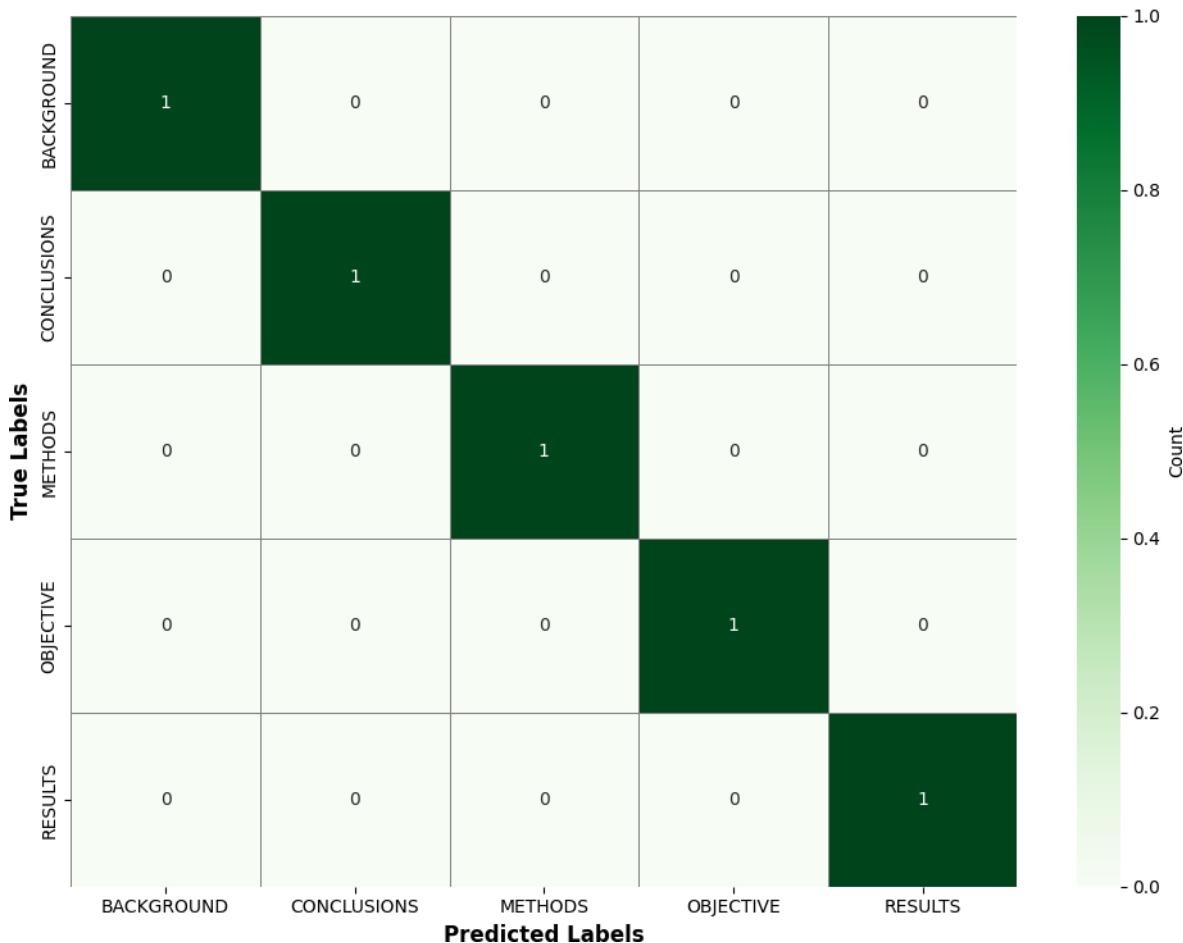
```
-----  
NaiveBayes : Accuracy = 1.0000 (100.00%)  
LogisticRegression : Accuracy = 1.0000 (100.00%)  
RandomForest : Accuracy = 1.0000 (100.00%)  
DecisionTree : Accuracy = 1.0000 (100.00%)  
KNN : Accuracy = 1.0000 (100.00%)  
-----
```

BOC (Hard Voting) Accuracy: 1.0000 (100.00%)  
BOC (Hard Voting) Macro F1 Score: 1.0000

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| BACKGROUND   | 1.00      | 1.00   | 1.00     | 1       |
| CONCLUSIONS  | 1.00      | 1.00   | 1.00     | 1       |
| METHODS      | 1.00      | 1.00   | 1.00     | 1       |
| OBJECTIVE    | 1.00      | 1.00   | 1.00     | 1       |
| RESULTS      | 1.00      | 1.00   | 1.00     | 1       |
| accuracy     |           |        | 1.00     | 5       |
| macro avg    | 1.00      | 1.00   | 1.00     | 5       |
| weighted avg | 1.00      | 1.00   | 1.00     | 5       |

**Confusion Matrix - Hard Voting (Bayes Optimal Approximation)**



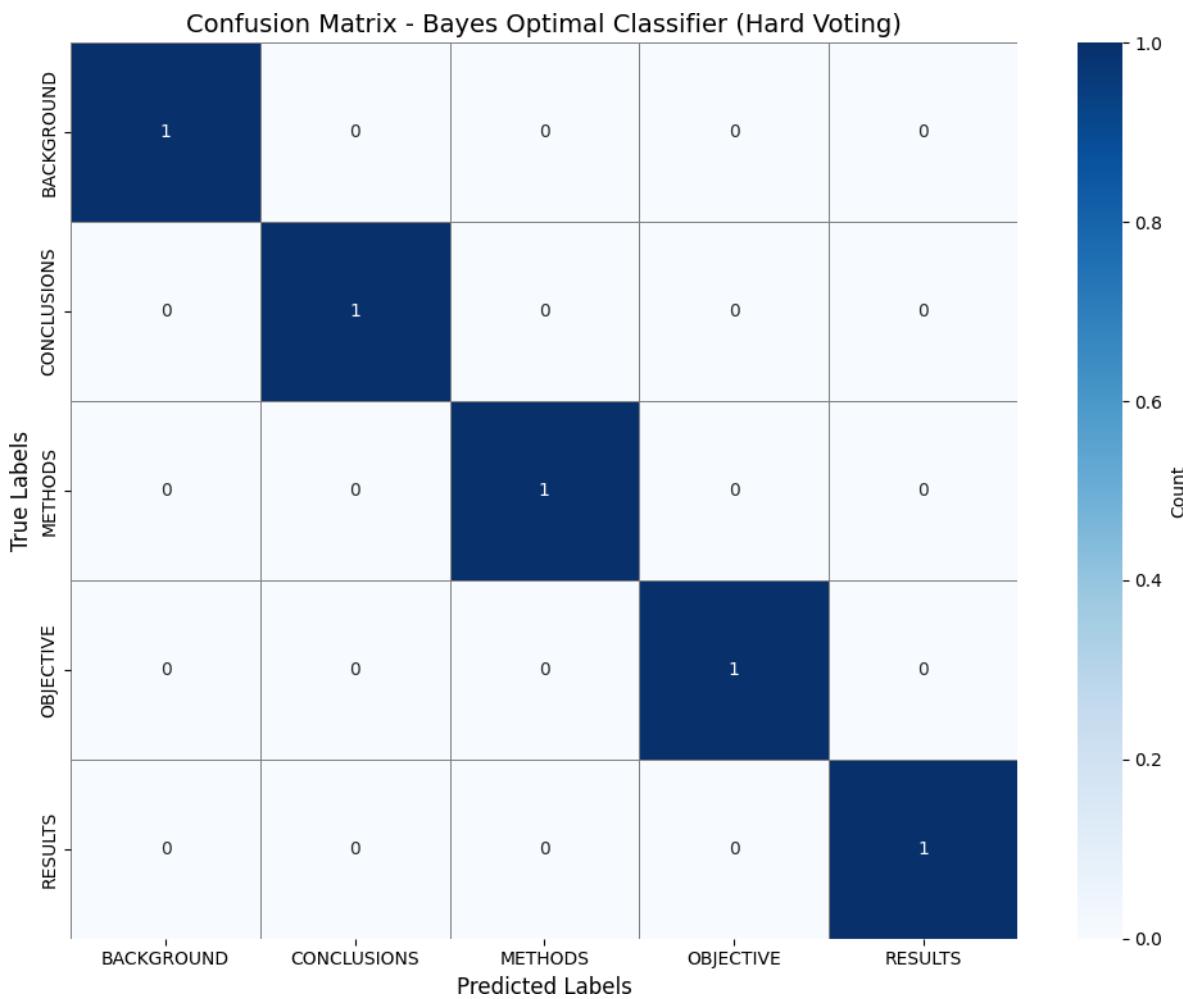
```

[+] === Final Evaluation: Bayes Optimal Classifier (Hard Voting) ===
BOC Accuracy: 1.0000
BOC Macro F1 Score: 1.0000

Classification Report:
precision    recall    f1-score   support
BACKGROUND    1.00    1.00    1.00      1
CONCLUSIONS   1.00    1.00    1.00      1
METHODS       1.00    1.00    1.00      1
OBJECTIVE     1.00    1.00    1.00      1
RESULTS        1.00    1.00    1.00      1

accuracy          1.00      5
macro avg         1.00    1.00    1.00      5
weighted avg      1.00    1.00    1.00      5

Confusion Matrix:
[[1 0 0 0 0]
 [0 1 0 0 0]
 [0 0 1 0 0]
 [0 0 0 1 0]
 [0 0 0 0 1]]
```



## Discussion

The **scratch model (Part A)** gave a decent performance but was limited because it used only basic word counts and simple probability estimation. The **tuned Sklearn model (Part B)** performed better after using TF-IDF features and optimized parameters, which improved accuracy and F1 score. Finally, the **BOC model (Part C)** achieved the best performance overall, since it combined multiple algorithms and used their collective strength. This ensemble approach reduced errors and provided more stable predictions compared to individual models.